

Introduction to Data Science

8 - Scikit Learn Basics



This Lecture

- Basic usage of the Scikit learn package

Setup

We'll typically start a notebook from now on with a set of standard imports, and any relevant Jupyter notebook "magic" directives:

```
In [1]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 import sklearn as sk
        4
        5 %matplotlib inline
```

Example Problem Setup

We'll continue to use our synthetic problem to illustrate linear regression in Scikit learn:

$$f(x) = 3 + 0.5n - n^2 + 0.15n^3$$
$$y = f(x) + \epsilon$$

```
In [2]: 1 # "ground truth" function
        2 def f(x):
        3     return 3 + 0.5 * x - x**2 + 0.15 * x**3
        4
```

Some more functions:

```
In [3]: 1 # ensure repeatability of this notebook
2 # (comment out for new results each run)
3 np.random.seed(12345)
4
5 # convenience function for generating samples
6 def sample(n, fn, limits, noise=1):
7     width = limits[1] - limits[0]
8     x = np.random.random(n) * width + limits[0]
9     y = fn(x) + np.random.randn(n) * noise
10    return x, y
11
12 # there's a scikit learn tool for generating
13 # polynomial features - this will be more useful
14 # when working with multivariate inputs, so we'll
15 # stick with this simpler solution for now
16 def phi(x, k):
17     return np.array([x ** p for p in range(k)]).T
```

Regression Workflow

- Obtain training samples (data and target)
- [optional] Do some initial visualization, statistics
- [optional] Preprocess data (generate features, dimensionality reduction)
- Initialize a model object
- Split data into training and test sets (or use cross validation, more on this another time)
- Train the model
- Use the trained model to make predictions
- Evaluate approximation quality (e.g., examine MSE)
- Visualize results
- Repeat steps as needed to refine model

Obtain Training Samples

Usually this involves getting data from an external source: the internet, your own research, etc.


For today, we'll simply sample from our synthetic problem.

```
In [4]: 1 n = 40
2
3 # we'll start using scikit learn's names for things
4 data, target = sample(n, f, [-5,5], 3)
```

Initial Visualization

This varies a lot. One common visualization is scatter plots showing correlations between pairs of inputs and/or the training data:

```
In [5]: 1 plt.scatter(data, target)
        2 plt.show()
```



Preprocess Data

Our initial visualization suggests non-linearity. Let's use some polynomial features.

```
In [6]: 1 Phi = phi(data, 5)
        2 Phi.shape
```

```
Out[6]: (40, 5)
```

Initialize a Model Object

Now we get to some actual Scikit learn code.

There are a bunch of regression models; we're going to use the [linear_model.LinearRegression](http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html#sklearn.linear_model.LinearRegression) (http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html#sklearn.linear_model.LinearRegression) one.

The basic process of obtaining a model, training, and then using it for prediction is uniform-ish across learning methods. Yay!

```
In [7]: 1 from sklearn.linear_model import LinearRegression
        2
        3 lr = LinearRegression(fit_intercept=False)
```

The `fit_intercept` parameter above defaults to True, but we already generated an intercept term in our design matrix.

Split Data into Training/Test Sets

There's a Scikit learn [function](http://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html#sklearn.model_selection.train_test_split) (http://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html#sklearn.model_selection.train_test_split) for this:

```
In [8]: 1 from sklearn.model_selection import train_test_split
2
3 X_train, X_test, y_train, y_test = train_test_split(
4     Phi, target, test_size = 0.5)
```

The `test_size` parameter is used to control what percentage of the data to hold out for testing.


We can check to see that our data matrices/vectors are the size expected:

```
In [9]: 1 print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(20, 5) (20, 5) (20,) (20,)
```

```
In [10]: 1 plt.plot(X_train[:,1], y_train, 'k.', X_test[:,1], y_test, 'r.')
```

```
Out[10]: [<matplotlib.lines.Line2D at 0x1bc496843d0>,
<matplotlib.lines.Line2D at 0x1bc49684370>]
```



Train the Model

This is done using the `fit` method of the model object:


```
In [11]: 1 lr.fit(X_train, y_train)
```

```
Out[11]: LinearRegression(fit_intercept=False)
```

Predict (and Visualize)


Let's see what we got:

```
In [12]: 1 x = np.arange(-5, 5, 0.1)
2 yhat = lr.predict(phi(x, 5))
3 plt.plot(x, yhat, 'r-', data, target, 'k.')
4 plt.show()
```



In our case, we also know ground truth, so let's add that in:

```
In [13]: 1 plt.plot(x, f(x), 'b-', x, yhat, 'r-', x_train[:,1], y_train, 'k.')
2 plt.show()
```



Evaluate

Let's compute MSE and RMSE on our test set:

```
In [14]: 1 MSE = ((y_test - lr.predict(X_test)) ** 2).mean()
2 RMSE = np.sqrt(MSE)
3 print("MSE: ", MSE)
4 print("RMSE:", RMSE)
```

```
MSE: 30.201240859904964
RMSE: 5.495565563243238
```

Refine Model

A lot we could do here! For now, let's repeat our work on evaluating RMSE for different orders.

I have to redo some work from above to get all the powers up to 12...

```
In [15]: 1 Phi = phi(data, 12)
2 X_train, X_test, y_train, y_test = train_test_split(
3     Phi, target, test_size = 0.5)
```

Now I can just pare down my feature matrices using NumPy's array slicing capabilities.


This works somewhat like list slicing, but with multi-dimensional support:

```
In [16]: 1 print(Phi.shape)
2
3 print(Phi[:, :3].shape)
```

```
(40, 12)
(40, 3)
```

Compute RMSEs across a range of orders and find the minimum:

```
In [18]: 1 RMSEs = []
2 orders = range(2,13)
3 for p in orders:
4     lr.fit(X_train[:,:(p+1)], y_train)
5     MSE = ((y_test - lr.predict(X_test[:,:(p+1)])) ** 2).mean()
6     RMSEs.append(np.sqrt(MSE))
7 RMSEs = np.array(RMSEs)
8 plt.plot(orders, RMSEs, 'b-x')
9 plt.plot(orders[RMSEs.argmin()], RMSEs.min(), 'rs') # marks the lowest RMSE value
10 plt.show()
11 print(RMSEs.min(), "at order", orders[RMSEs.argmin()])
```



3.1110201095243624 at order 3

```
In [19]: 1 lr.fit(X_train[:, :5], y_train)
```

```
Out[19]: LinearRegression(fit_intercept=False)
```


```
In [20]: 1 lr.score(X_test[:,5], y_test)
```

```
Out[20]: 0.8768295733190168
```

```
In [21]: 1 X_test.shape
```

```
Out[21]: (20, 12)
```

```
In [23]: 1 COEFFs = []  
2 for p in orders:  
3     lr.fit(X_train[:,p], y_train)  
4     COEFFs.append(lr.score(X_test[:,p], y_test))  
5 COEFFs = np.array(COEFFs)  
6 plt.plot(orders, COEFFs, 'b-x')  
7 plt.plot(orders[COEFFs.argmax()], COEFFs.max(), 'rs') # plots the max value with a  
8 plt.show()  
9 print(COEFFs.max(), "at order", orders[COEFFs.argmax()])
```



```
0.8833925877665701 at order 4
```

```
In [ ]: 1
```