

Introduction to Data Science

11 - Data Acquisition

This Lecture

- Importing data using Python and pandas

The obligatory setup code...

```
In [2]: 1 import numpy as np
        2 import pandas as pd
        3
        4 from pandas import Series, DataFrame
```

Raw Python

Text-based files are easy to read and write in Python.

In particular, files which data is organized into individual lines.

There's a sample file, `text.txt`, located in the directory `11-acquisition-files`.

We can view it by asking Jupyter to execute the linux command `cat` on the file (only works if you are not using Windows):

```
In [3]: 1 !cat 11-acquisition-files/text.txt
```

| Name | Age | Salary | Hired |
|--------|-----|---------|-----------|
| Laura | 52 | 103,790 | 1/1/2005 |
| Shashi | 46 | 89,100 | 6/16/2010 |
| Jun | 33 | 85,500 | 7/1/2017 |
| Bruce | 48 | 96,445 | 12/1/2008 |
| Raluca | 40 | 110,080 | 9/15/2012 |

Reading the file via Python is pretty simple (use this if you are using a Windows operating system):

```
In [4]: 1 f = open('11-acquisition-files/text.txt') # get an open file object
        2 for line in f:                          # for loops work!
        3     print(line, end='')                  # print without an newline
        4
        5 f.close()
```

| Name | Age | Salary | Hired |
|--------|-----|---------|-----------|
| Laura | 52 | 103,790 | 1/1/2005 |
| Shashi | 46 | 89,100 | 6/16/2010 |
| Jun | 33 | 85,500 | 7/1/2017 |
| Bruce | 48 | 96,445 | 12/1/2008 |
| Raluca | 40 | 110,080 | 9/15/2012 |

Python has many tools to let us relatively easily parse files like this.

For simple files, the string `split` method may suffice:

```
In [5]: 1 f = open('11-acquisition-files/text.txt')
        2 for line in f:
        3     print(line.split())
        4 f.close()
```

```
['Name', 'Age', 'Salary', 'Hired']
['Laura', '52', '103,790', '1/1/2005']
['Shashi', '46', '89,100', '6/16/2010']
['Jun', '33', '85,500', '7/1/2017']
['Bruce', '48', '96,445', '12/1/2008']
['Raluca', '40', '110,080', '9/15/2012']
```

If you know regular expressions, you can parse more complicated files.

However, you're unlikely to need something that complicated.

Let's bring this data into pandas, the hard way:

```
In [6]: 1 data = []
        2 f = open('11-acquisition-files/text.txt')
        3 columns = next(f).split()
        4 for line in f:
        5     data.append(line.split())
        6 f.close()
        7 df = DataFrame(data=data, columns=columns)
        8 df
```

Out[6]:

| | Name | Age | Salary | Hired |
|---|--------|-----|---------|-----------|
| 0 | Laura | 52 | 103,790 | 1/1/2005 |
| 1 | Shashi | 46 | 89,100 | 6/16/2010 |
| 2 | Jun | 33 | 85,500 | 7/1/2017 |
| 3 | Bruce | 48 | 96,445 | 12/1/2008 |
| 4 | Raluca | 40 | 110,080 | 9/15/2012 |

pandas read_table

For many file formats, it is easiest to let pandas do all the work!

One of the most basic tools is used to read text files like the one we worked with above is `read_table`.

```
In [7]: 1 pd.read_table?
```

We don't need to do much for this file, as it is nicely tab-delimited and has a header row.

```
In [8]: 1 df2 = pd.read_table('11-acquisition-files/text.txt')
2 df2
```

Out[8]:

| | Name | Age | Salary | Hired |
|---|--------|-----|---------|-----------|
| 0 | Laura | 52 | 103,790 | 1/1/2005 |
| 1 | Shashi | 46 | 89,100 | 6/16/2010 |
| 2 | Jun | 33 | 85,500 | 7/1/2017 |
| 3 | Bruce | 48 | 96,445 | 12/1/2008 |
| 4 | Raluca | 40 | 110,080 | 9/15/2012 |

Our table might not have a header row, in which case we have to supply column labels:

```
In [9]: 1 df3 = pd.read_table('11-acquisition-files/text-no-header.txt',
2                             header=None, names=['Person', 'Age', 'Income'])
3 df3
```

Out[9]:

| | Person | Age | Income |
|---|--------|-----|---------|
| 0 | Laura | 52 | 103,790 |
| 1 | Shashi | 46 | 89,100 |
| 2 | Jun | 33 | 85,500 |
| 3 | Bruce | 48 | 96,445 |
| 4 | Raluca | 40 | 110,080 |

There are lots of other ways to manipulate the data when reading.

For instance, we can make one column the row index:

```
In [10]: 1 df4 = pd.read_table('11-acquisition-files/text.txt',
2                             index_col=0)
3 df4
```

Out[10]:

| | Age | Salary | Hired |
|---------------|-----|---------|-----------|
| Name | | | |
| Laura | 52 | 103,790 | 1/1/2005 |
| Shashi | 46 | 89,100 | 6/16/2010 |
| Jun | 33 | 85,500 | 7/1/2017 |
| Bruce | 48 | 96,445 | 12/1/2008 |
| Raluca | 40 | 110,080 | 9/15/2012 |

We can also do this by giving the column header:

```
In [11]: 1 df5 = pd.read_table('11-acquisition-files/text.txt',
2                             index_col='Name')
3 df5
```

Out[11]:

| | Age | Salary | Hired |
|---------------|-----|---------|-----------|
| Name | | | |
| Laura | 52 | 103,790 | 1/1/2005 |
| Shashi | 46 | 89,100 | 6/16/2010 |
| Jun | 33 | 85,500 | 7/1/2017 |
| Bruce | 48 | 96,445 | 12/1/2008 |
| Raluca | 40 | 110,080 | 9/15/2012 |

If tabs are not your separator, you can specify separators either as a specific character, or using a regular expression.

For arbitrary length whitespace separation, use `'\s+'` :

```
In [12]: 1 df6 = pd.read_table('11-acquisition-files/text.txt', sep='\s+')
2 df6
```

Out[12]:

| | Name | Age | Salary | Hired |
|----------|--------|-----|---------|-----------|
| 0 | Laura | 52 | 103,790 | 1/1/2005 |
| 1 | Shashi | 46 | 89,100 | 6/16/2010 |
| 2 | Jun | 33 | 85,500 | 7/1/2017 |
| 3 | Bruce | 48 | 96,445 | 12/1/2008 |
| 4 | Raluca | 40 | 110,080 | 9/15/2012 |

Don't want the whole thing?

Check out the `nrows` and `skiprows` arguments:

```
In [13]: 1 df7 = pd.read_table('11-acquisition-files/text.txt', nrows=3)
         2 df7
```

Out[13]:

| | Name | Age | Salary | Hired |
|---|--------|-----|---------|-----------|
| 0 | Laura | 52 | 103,790 | 1/1/2005 |
| 1 | Shashi | 46 | 89,100 | 6/16/2010 |
| 2 | Jun | 33 | 85,500 | 7/1/2017 |

```
In [14]: 1 df8 = pd.read_table('11-acquisition-files/text.txt',
         2                      skiprows=[0,1])
         3 df8
```

Out[14]:

| | Shashi | 46 | 89,100 | 6/16/2010 |
|---|--------|----|---------|-----------|
| 0 | Jun | 33 | 85,500 | 7/1/2017 |
| 1 | Bruce | 48 | 96,445 | 12/1/2008 |
| 2 | Raluca | 40 | 110,080 | 9/15/2012 |

Type Inference in `read_table`

Let's take a closer look at the DataFrame objects we're getting back:

```
In [15]: 1 df9 = pd.read_table('11-acquisition-files/text.txt')
         2 print(df9.info())
         3 df9
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 4 columns):
Name      5 non-null object
Age       5 non-null int64
Salary    5 non-null object
Hired     5 non-null object
dtypes: int64(1), object(3)
memory usage: 240.0+ bytes
None
```

Out[15]:

| | Name | Age | Salary | Hired |
|---|--------|-----|---------|-----------|
| 0 | Laura | 52 | 103,790 | 1/1/2005 |
| 1 | Shashi | 46 | 89,100 | 6/16/2010 |
| 2 | Jun | 33 | 85,500 | 7/1/2017 |
| 3 | Bruce | 48 | 96,445 | 12/1/2008 |
| 4 | Raluca | 40 | 110,080 | 9/15/2012 |

It correctly *inferred* the type of the Age column as an integer.

But it didn't pick up Salary as a number, which we kind of want.

OK, this can be fixed!

```
In [16]: 1 df10 = pd.read_table('11-acquisition-files/text.txt',
2                             thousands=',')
3 df10.info()
4 #pd.read_table?
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 4 columns):
Name      5 non-null object
Age       5 non-null int64
Salary    5 non-null int64
Hired     5 non-null object
dtypes: int64(2), object(2)
memory usage: 240.0+ bytes
```

Dates are also something pandas knows about:

```
In [17]: 1 df11 = pd.read_table('11-acquisition-files/text.txt',
2                             thousands=',',
3                             parse_dates=['Hired'])
4 df11.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 4 columns):
Name      5 non-null object
Age       5 non-null int64
Salary    5 non-null int64
Hired     5 non-null datetime64[ns]
dtypes: datetime64[ns](1), int64(2), object(1)
memory usage: 240.0+ bytes
```

The display data now looks slightly different:

```
In [18]: 1 df11
```

Out[18]:

| | Name | Age | Salary | Hired |
|---|--------|-----|--------|------------|
| 0 | Laura | 52 | 103790 | 2005-01-01 |
| 1 | Shashi | 46 | 89100 | 2010-06-16 |
| 2 | Jun | 33 | 85500 | 2017-07-01 |
| 3 | Bruce | 48 | 96445 | 2008-12-01 |
| 4 | Raluca | 40 | 110080 | 2012-09-15 |

pandas read_csv

A startling amount of the time, data is available in *comma-separated values* format.

The CSV format is commonly used to exchange data between things like spreadsheets and databases.

Here's what our data might look like exported from a spreadsheet program:

```
In [19]: 1 !cat ll-acquisition-files/text.csv
```

```
Name, Age, Salary, Hired
Laura, 52, 103790, 1/1/2005
Shashi, 46, 89100, 6/16/2010
Jun, 33, 85500, 7/1/2017
Bruce, 48, 96445, 12/1/2008
Raluca, 40, 110080, 9/15/2012
```

pandas `read_csv` is pretty much identical to `read_table` , but assumes a CSV format:

```
In [20]: 1 df12 = pd.read_csv('ll-acquisition-files/text.csv',
2                        parse_dates=[ 'Hired' ])
3 df12
```

Out[20]:

| | Name | Age | Salary | Hired |
|---|--------|-----|--------|------------|
| 0 | Laura | 52 | 103790 | 2005-01-01 |
| 1 | Shashi | 46 | 89100 | 2010-06-16 |
| 2 | Jun | 33 | 85500 | 2017-07-01 |
| 3 | Bruce | 48 | 96445 | 2008-12-01 |
| 4 | Raluca | 40 | 110080 | 2012-09-15 |

JSON Data

Increasingly data is available in JSON (*JavaScript Object Notation*) format.

JSON is a format for exchanging rich, structured data as plain text.

JSON object strings look (and act) remarkably like Python code for dictionaries:

```
{ "Name" : "Laura", "Age" : 52, "Salary" : 103790, "Pets" : [ { "type" : "rabbit", "name" : "Gandalf" }, { "type" : "dog", "name" : "Aragorn" } ] }
```

JSON basic types include strings, numbers (integer or floating point), Booleans, and nulls.

Compound types are objects and lists, which correspond pretty directly to Python dicts and lists.

Python has a library for interpreting JSON strings and turning them into Python objects (and vice versa):

```
In [21]: 1 import json
2
3 s = '''
4     { "Name" : "Laura",
5       "Age" : 52,
6       "Salary" : 103790,
7       "Pets" : [
8         { "type" : "rabbit", "name" : "Gandalf" },
9         { "type" : "dog", "name" : "Aragorn" } ] }
10     '''
11
12 obj = json.loads(s)
13 obj
```

```
Out[21]: {'Name': 'Laura',
'Age': 52,
'Salary': 103790,
'Pets': [{'type': 'rabbit', 'name': 'Gandalf'},
{'type': 'dog', 'name': 'Aragorn'}]}
```

It turns out that pandas DataFrame objects can be created from Python dictionaries, although its approach isn't always exactly what you expect:

```
In [22]: 1 df13 = DataFrame(obj)
2 df13
```

```
Out[22]:
```

| | Name | Age | Salary | Pets |
|---|-------|-----|--------|---------------------------------------|
| 0 | Laura | 52 | 103790 | {'type': 'rabbit', 'name': 'Gandalf'} |
| 1 | Laura | 52 | 103790 | {'type': 'dog', 'name': 'Aragorn'} |

When dictionaries are fairly "flat", pandas interprets them pretty sensibly.

The two structures it handles best are dictionaries where the keys represent columns, and lists, where the entries are dictionaries representing rows.

Here's an example for the first:

```
In [23]: 1 d = { "Name" : [ "Laura", "Shashi", "Jun", "Bruce", "Raluca" ],
2         "Age" : [ 52, 46, 33, 48, 40 ],
3         "Salary" : [ 103790, 89100, 85500, 96445, 110080 ],
4         "Hired" : [ "1/1/2005", "6/16/2010", "7/1/2017", "12/1/2008", "9/15/2012" ] }
5 DataFrame(d)
```

```
Out[23]:
```

| | Name | Age | Salary | Hired |
|---|--------|-----|--------|-----------|
| 0 | Laura | 52 | 103790 | 1/1/2005 |
| 1 | Shashi | 46 | 89100 | 6/16/2010 |
| 2 | Jun | 33 | 85500 | 7/1/2017 |
| 3 | Bruce | 48 | 96445 | 12/1/2008 |
| 4 | Raluca | 40 | 110080 | 9/15/2012 |

and the second:


```
In [24]: 1 d2 = [
2         {"Name" : "Laura", "Age" : 52, "Salary" : 103790, "Hired" : "1/1/2005"},
3         {"Name" : "Shashi", "Age" : 46, "Salary" : 89100, "Hired" : "6/16/2010"},
4         {"Name" : "Jun", "Age" : 33, "Salary" : 85500, "Hired" : "7/1/2017"},
5         {"Name" : "Bruce", "Age" : 48, "Salary" : 96445, "Hired" : "12/1/2008"},
6         {"Name" : "Raluca", "Age" : 40, "Salary" : 110080, "Hired" : "9/15/2012"}
7     ]
8     DataFrame(d2)
```

Out[24]:

| | Age | Hired | Name | Salary |
|---|-----|-----------|--------|--------|
| 0 | 52 | 1/1/2005 | Laura | 103790 |
| 1 | 46 | 6/16/2010 | Shashi | 89100 |
| 2 | 33 | 7/1/2017 | Jun | 85500 |
| 3 | 48 | 12/1/2008 | Bruce | 96445 |
| 4 | 40 | 9/15/2012 | Raluca | 110080 |

In both cases, pandas re-ordered the columns alphabetically; we can tell it what order we want things in:

```
In [25]: 1 DataFrame(d2, columns=['Name', 'Age', 'Salary', 'Hired'])
```

Out[25]:

| | Name | Age | Salary | Hired |
|---|--------|-----|--------|-----------|
| 0 | Laura | 52 | 103790 | 1/1/2005 |
| 1 | Shashi | 46 | 89100 | 6/16/2010 |
| 2 | Jun | 33 | 85500 | 7/1/2017 |
| 3 | Bruce | 48 | 96445 | 12/1/2008 |
| 4 | Raluca | 40 | 110080 | 9/15/2012 |

Additional structures are possible, such as ones in which row index labels are explicitly provided:

```
In [26]: 1 d3 = {"Age" : {"Laura" : 52, "Shashi" : 46},
2         "Salary" : {"Laura" : 103790, "Shashi" : 89100},
3         "Hired" : {"Laura" : "1/1/2005", "Shashi" : "6/16/2010"}}
4     DataFrame(d3)
```

Out[26]:

| | Age | Salary | Hired |
|--------|-----|--------|-----------|
| Laura | 52 | 103790 | 1/1/2005 |
| Shashi | 46 | 89100 | 6/16/2010 |

Not surprisingly, then, if data is stored in JSON in any of these structures, pandas can read it quite easily.

Here is our simple data set in JSON data formats:

```
In [27]: 1 !cat ll-acquisition-files/json1.json
2 d14 = pd.read_json('ll-acquisition-files/json1.json')
3 d14
```

```
[{"Name" : "Laura", "Age" : 52, "Salary" : 103790, "Hired" : "1/1/2005"},
{"Name" : "Shashi", "Age" : 46, "Salary" : 89100, "Hired" : "6/16/2010"},
{"Name" : "Jun", "Age" : 33, "Salary" : 85500, "Hired" : "7/1/2017"},
{"Name" : "Bruce", "Age" : 48, "Salary" : 96445, "Hired" : "12/1/2008"},
{"Name" : "Raluca", "Age" : 40, "Salary" : 110080, "Hired" : "9/15/2012"}]
```

Out[27]:

| | Age | Hired | Name | Salary |
|---|-----|-----------|--------|--------|
| 0 | 52 | 1/1/2005 | Laura | 103790 |
| 1 | 46 | 6/16/2010 | Shashi | 89100 |
| 2 | 33 | 7/1/2017 | Jun | 85500 |
| 3 | 48 | 12/1/2008 | Bruce | 96445 |
| 4 | 40 | 9/15/2012 | Raluca | 110080 |

```
In [28]: 1 !cat ll-acquisition-files/json2.json
2 d15 = pd.read_json('ll-acquisition-files/json2.json',
3                   convert_dates=['Hired'])
4 d15
```

```
{ "Name"      : ["Laura", "Shashi", "Jun", "Bruce", "Raluca"],
  "Age"       : [52, 46, 33, 48, 40],
  "Salary"    : [103790, 89100, 85500, 96445, 110080],
  "Hired"     : ["1/1/2005", "6/16/2010", "7/1/2017", "12/1/2008", "9/15/2012"] }
```

Out[28]:

| | Name | Age | Salary | Hired |
|---|--------|-----|--------|------------|
| 0 | Laura | 52 | 103790 | 2005-01-01 |
| 1 | Shashi | 46 | 89100 | 2010-06-16 |
| 2 | Jun | 33 | 85500 | 2017-07-01 |
| 3 | Bruce | 48 | 96445 | 2008-12-01 |
| 4 | Raluca | 40 | 110080 | 2012-09-15 |

Oddly, there doesn't seem to be a way to set the order of columns in `read_json` , but we can reorder pretty easily:

```
In [29]: 1 d16 = d15[['Name', 'Age', 'Salary', 'Hired']]
          2 d16
```

Out[29]:

| | Name | Age | Salary | Hired |
|----------|-------------|------------|---------------|--------------|
| 0 | Laura | 52 | 103790 | 2005-01-01 |
| 1 | Shashi | 46 | 89100 | 2010-06-16 |
| 2 | Jun | 33 | 85500 | 2017-07-01 |
| 3 | Bruce | 48 | 96445 | 2008-12-01 |
| 4 | Raluca | 40 | 110080 | 2012-09-15 |