


Introduction to Data Science

18 - Unsupervised Learning



This Lecture

- Introduction to unsupervised learning
- Data preprocessing
 - Scaling and normalization
 - Dimensionality reduction

Setup

The obligatory setup code.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import sklearn as sk
        4 import matplotlib.pyplot as plt
        5
        6 from pandas import DataFrame
        7
        8 plt.style.use("ggplot")
        9
       10 %matplotlib inline
       11 %config InlineBackend.figure_format = 'retina'
```

```
In [2]: 1 # function for generating normally distributed data/
2 def sample_cluster(n, x, y, sigma):
3     x = np.random.randn(n) * sigma + x;
4     y = np.random.randn(n) * sigma + y;
5     return np.array([x, y]).T
6
```


Unsupervised vs Supervised

In supervised learning, we have *labeled* data:

- some input variables
- some additional variable(s) which we are learning to predict


For example, we might have a classification problem like the one below (colors = class labels):

```
In [4]: 1 np.random.seed(1234)
2
3 # creates three 'sets' of randomly distributed data
4 c1 = sample_cluster(50, 0, 0.5, 0.15)
5 c2 = sample_cluster(50, 0, 1.2, 0.1)
6 c3 = sample_cluster(50, 0.5, 1, 0.2)
7
8 # plots those three sets (c1 is red. c2 is blue. c3 is yellow)
9 plt.plot(c1[:,0], c1[:,1], 'r.', c2[:,0], c2[:,1], 'b.', c3[:,0], c3[:,1], 'y.',)
10 plt.show()
```



In *unsupervised* learning, we are given no labels, and we seek to find hidden patterns in the data:

```
In [6]: 1 # since there are no labels we want the data to look the same (make them all black
2 plt.plot(c1[:,0], c1[:,1], 'k.', c2[:,0], c2[:,1], 'k.', c3[:,0], c3[:,1], 'k.',)
3 plt.show()
```



Questions we could ask about the data:

- Is there a transformation of the data which will reveal patterns (to humans or algorithms)?
- What are the relevant features of the data which are informative?
- Are there natural groupings into which we could separate the data?

Challenges of Unsupervised Learning

Since we have no labeled data, there are no predictions that we can make *and meaningfully test*.

Evaluation of unsupervised learning algorithms is often largely subjective.

Unsupervised learning is often used in *exploratory data analysis*.

Example Applications

- Group (cluster) gene expression data in cancer patients to look for patterns; a gene (or group of genes) which strongly differentiates patients may be worth further study:
 - Different disease causes
 - Different responses to treatment
- Anomaly Detection: E.g., look for anomalous patterns in credit card spending
- Group people or organizations according to some new identifiers

- Reveal hidden similarities
- Provide alerts to activities with similar risks (e.g., fund analysis)
- Targeted marketing

Data Preprocessing

- Generally useful to improve supervised learning algorithm performance
- Scaling/normalization:
 - Transform data so that features are on same scale or have same statistics
 - Helps some algorithms which are sensitive to scale
- Dimensionality reduction:
 - Transform data into a **sub-space in which visualization and/or machine learning is easier**
 - Reduce computational cost of learning

Scaling

Is a thing. It helps with some machine learning algorithms.

Two common ones with numpy examples for a vector of values for a single feature, `x` :

Standard scaling: `x_scaled = (x - np.mean(x)) / np.std(x)`

Min-max scaling (to [0, 1] range): `x_scaled = (x - np.min(x)) / (np.max(x) - np.min(x))`

Dimensionality Reduction

Problem:

- Input data is often has (very) high dimensionality (100s to 1000s)
- This can lead to expensive learning and promotes overfitting
- Variables can often also have high correlation (redundant information, but more noise, again promoting overfitting)

Solution:

- Extract most relevant sub-space of input data before visualizing or using machine learning

Dimensionality reduction can be used as preprocessing for either supervised or unsupervised learning

Principal Components Analysis

The most popular form of dimensionality reduction.

Lots of linear algebra behind this. We won't go there.

Basically, rotates and transforms the data into a **new, high-dimensional space**.


Coordinate system of the new space is such that:

- Dimensions (axes, features) are ordered by the amount of **"fraction of variance explained"** (relevance)
- Feature values along 1st dimension give highest variance explained
- Feature values along 2nd dimension give 2nd highest variance explained
- Etc.

A subspace (fewer dimensions) of the original data space can be created by discarding dimensions that have low variance explained.


A dimension with a high variance explained may indicate that two or more distinct groups fall on different locations along that dimension.
E.g....

```
In [7]: 1 # Code is only for creating figure below. Not instructive of how to execute PCA.
2 np.random.seed(1234)
3 c1 = sample_cluster(50, 1, 2, 0.1)
4 c2 = sample_cluster(50, 2, 4, 0.1)
5 plt.figure(figsize=(8, 8))
6 plt.plot(c1[:,0], c1[:,1], 'r.', c2[:,0], c2[:,1], 'b.')
7 ax = plt.gca()
8 ax.axes.set_aspect('equal')
9 plt.axis([-5, 5, -5, 5])
10 ax1 = plt.plot([-2.5, 2.5], [-5, 5], 'c-', label='PC axis 1')
11 ax2 = plt.plot([-5, 5], [2.5, -2.5], 'm-', label='PC axis 2')
12 plt.legend(fontsize=16)
13 plt.show()
```



More commonly, the data are simply "spread" more widely along a dimension of high variance explained than along a dimension of low variance explained.
E.g....

```
In [8]: 1 # Code is only for creating figure below. Not instructive of how to execute PCA.
2 np.random.seed(1234)
3 x = np.random.randn(50)*2
4 y = 0.3 * x + np.random.randn(50)/5
5 plt.figure(figsize=(8, 8))
6 plt.plot(x, y, 'b.')
7 plt.plot()
8 ax = plt.gca()
9 ax.axes.set_aspect('equal')
10 plt.axis([-5, 5, -5, 5])
11 s = 1.5
12 ax1 = plt.plot([-5, 5], [-s, s], 'c-', label='PC axis 1')
13 ax2 = plt.plot([-s, s], [5, -5], 'm-', label='PC axis 2')
14 plt.legend(fontsize=16)
15 plt.show()
```




PCA dimensionality reduction example

Consider this dataset:


```
In [9]: 1 M = [[1, -1, 7],[20, 3, -5],[1,1,1]]
2
3 x1 = np.random.randn(300);
4 y1 = np.random.randn(300);
5 z1 = np.random.randn(300);
6 data1 = np.array([x1, y1, z1]).T @ M
7
8 x2 = np.random.randn(300);
9 y2 = np.random.randn(300);
10 z2 = np.random.randn(300);
11 data2 = np.array([x2, y2, z2]).T @ M + np.array([20,-10,15])
12
13 data = np.concatenate((data1, data2))
```

```
In [10]: 1 from mpl_toolkits.mplot3d import Axes3D
2 fig = plt.figure()
3 ax = fig.add_subplot(111, projection='3d')
4 ax.scatter(data1[:,0], data1[:,1], data1[:,2])
5 ax.scatter(data2[:,0], data2[:,1], data2[:,2])
6 plt.show()
```



Let's apply PCA and look at the first two principal components.

```
In [11]: 1 from sklearn.decomposition import PCA
2         pca = PCA(n_components=2)
3         pca.fit(data)
4         data1_pca = pca.transform(data1)
5         data2_pca = pca.transform(data2)
6
7         plt.scatter(data1_pca[:,0], data1_pca[:,1])
8         plt.scatter(data2_pca[:,0], data2_pca[:,1])
9         plt.show()
```



We've reduced the dimensionality from 3 to 2, while maintaining the separability of the two data classes (for the most part).

Taiwan Credit Card Default Dataset

- Real data aren't nearly as pretty

```
In [10]: 1 data = pd.read_csv('default.csv', header=1, encoding='utf8', index_col='ID')
2         data.head()
```

Out[10]:

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5	...	BILL_AMT4	BILL_
ID													
1	20000	2	2	1	24	2	2	-1	-1	-2	...	0	
2	120000	2	2	2	26	-1	2	0	0	0	...	3272	
3	90000	2	2	2	34	0	0	0	0	0	...	14331	
4	50000	2	2	1	37	0	0	0	0	0	...	28314	
5	50000	1	2	1	57	-1	0	-1	0	0	...	20940	

5 rows × 24 columns


```
In [11]: 1 # Use get_dummies to convert categorical features into sets of Boolean features (m
2 all_dummies = ['SEX', 'EDUCATION', 'MARRIAGE', 'PAY_0', 'PAY_2', 'PAY_3', 'PAY_4', 'PAY_5
3 df3 = pd.get_dummies(data, columns=all_dummies)
4 df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30000 entries, 1 to 30000
Data columns (total 92 columns):
LIMIT_BAL                30000 non-null int64
AGE                      30000 non-null int64
BILL_AMT1                30000 non-null int64
BILL_AMT2                30000 non-null int64
BILL_AMT3                30000 non-null int64
BILL_AMT4                30000 non-null int64
BILL_AMT5                30000 non-null int64
BILL_AMT6                30000 non-null int64
PAY_AMT1                 30000 non-null int64
PAY_AMT2                 30000 non-null int64
PAY_AMT3                 30000 non-null int64
PAY_AMT4                 30000 non-null int64
PAY_AMT5                 30000 non-null int64
PAY_AMT6                 30000 non-null int64
default payment next month 30000 non-null int64
SEX_1                    30000 non-null uint8
SEX_2                    30000 non-null uint8
EDUCATION_0              30000 non-null uint8
EDUCATION_1              30000 non-null uint8
EDUCATION_2              30000 non-null uint8
EDUCATION_3              30000 non-null uint8
EDUCATION_4              30000 non-null uint8
EDUCATION_5              30000 non-null uint8
EDUCATION_6              30000 non-null uint8
MARRIAGE_0               30000 non-null uint8
MARRIAGE_1               30000 non-null uint8
MARRIAGE_2               30000 non-null uint8
MARRIAGE_3               30000 non-null uint8
PAY_0_-2                 30000 non-null uint8
PAY_0_-1                 30000 non-null uint8
PAY_0_0                  30000 non-null uint8
PAY_0_1                  30000 non-null uint8
PAY_0_2                  30000 non-null uint8
PAY_0_3                  30000 non-null uint8
PAY_0_4                  30000 non-null uint8
PAY_0_5                  30000 non-null uint8
PAY_0_6                  30000 non-null uint8
PAY_0_7                  30000 non-null uint8
PAY_0_8                  30000 non-null uint8
PAY_2_-2                 30000 non-null uint8
PAY_2_-1                 30000 non-null uint8
PAY_2_0                  30000 non-null uint8
PAY_2_1                  30000 non-null uint8
PAY_2_2                  30000 non-null uint8
PAY_2_3                  30000 non-null uint8
PAY_2_4                  30000 non-null uint8
PAY_2_5                  30000 non-null uint8
PAY_2_6                  30000 non-null uint8
PAY_2_7                  30000 non-null uint8
PAY_2_8                  30000 non-null uint8
PAY_3_-2                 30000 non-null uint8
PAY_3_-1                 30000 non-null uint8
PAY_3_0                  30000 non-null uint8
PAY_3_1                  30000 non-null uint8
```

```

PAY_3_2          30000 non-null uint8
PAY_3_3          30000 non-null uint8
PAY_3_4          30000 non-null uint8
PAY_3_5          30000 non-null uint8
PAY_3_6          30000 non-null uint8
PAY_3_7          30000 non-null uint8
PAY_3_8          30000 non-null uint8
PAY_4_-2         30000 non-null uint8
PAY_4_-1         30000 non-null uint8
PAY_4_0          30000 non-null uint8
PAY_4_1          30000 non-null uint8
PAY_4_2          30000 non-null uint8
PAY_4_3          30000 non-null uint8
PAY_4_4          30000 non-null uint8
PAY_4_5          30000 non-null uint8
PAY_4_6          30000 non-null uint8
PAY_4_7          30000 non-null uint8
PAY_4_8          30000 non-null uint8
PAY_5_-2         30000 non-null uint8
PAY_5_-1         30000 non-null uint8
PAY_5_0          30000 non-null uint8
PAY_5_2          30000 non-null uint8
PAY_5_3          30000 non-null uint8
PAY_5_4          30000 non-null uint8
PAY_5_5          30000 non-null uint8
PAY_5_6          30000 non-null uint8
PAY_5_7          30000 non-null uint8
PAY_5_8          30000 non-null uint8
PAY_6_-2         30000 non-null uint8
PAY_6_-1         30000 non-null uint8
PAY_6_0          30000 non-null uint8
PAY_6_2          30000 non-null uint8
PAY_6_3          30000 non-null uint8
PAY_6_4          30000 non-null uint8
PAY_6_5          30000 non-null uint8
PAY_6_6          30000 non-null uint8
PAY_6_7          30000 non-null uint8
PAY_6_8          30000 non-null uint8
dtypes: int64(15), uint8(77)
memory usage: 5.9 MB

```

Pretend this is a supervised learning problem in which we want a model that predicts whether there will be a payment default in the next month.

```

In [12]: 1 # Extract separate DataFrames (or Series) for the predictors and the targets
          2 target = 'default payment next month'
          3 inputs3 = df3.columns.drop(target)
          4
          5 X = df3[inputs3]
          6 t = df3[target]
          7
          8 X.shape

```

Out[12]: (30000, 91)

We have a lot of predictors: 91 (though this isn't super high for a data set with lots of samples--30,000 in this case).

This could make a trained model prone to overfitting. Let's use PCA to lower the number of dimensions, and see if we learn anything by visualizing the two PC dimensions with the highest variance explained.


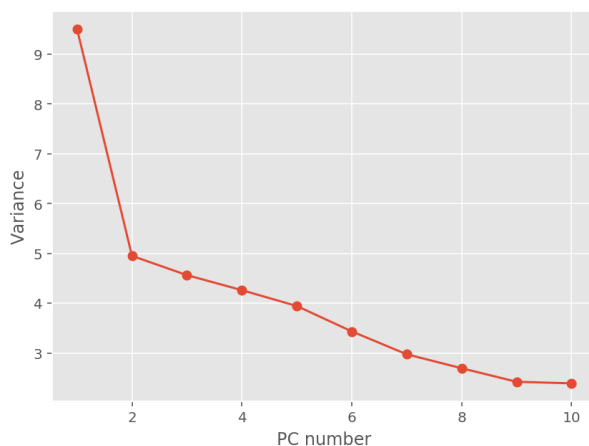
```
In [13]: 1 # First we'll apply standard scaling to the predictors.
2
3 from sklearn.preprocessing import StandardScaler
4
5 # Calculate the scaling parameters
6 ss = StandardScaler()
7 ss.fit(X)
8
9 # Perform the scaling
10 X_scaled = ss.transform(X)
11
12 #print(X)
13 #print(X_scaled)
```

```
In [14]: 1 # Now we'll apply PCA, and keep the top 10 components/dimensions
2
3 # Use PCA to compute the transformation parameters
4 pca = PCA(n_components=10)
5 pca.fit(X_scaled)
6
7 # Apply the transformation to our original data
8 X_pca = pca.transform(X_scaled)
```

```
In [15]: 1 X.shape, X_pca.shape
```

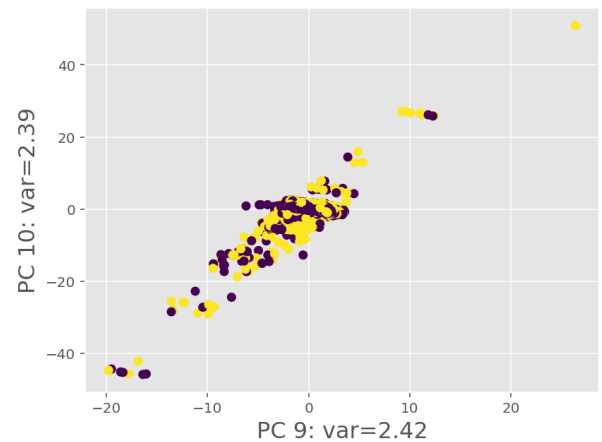

```
Out[15]: ((30000, 91), (30000, 10))
```

```
In [16]: 1 # Let's look at the variance (and variance explained) for each component
2
3 plt.figure(figsize=(15, 5))
4 plt.subplot(1,2,1)
5 plt.plot(np.arange(1,11), pca.explained_variance_, 'o-')
6 plt.xlabel('PC number')
7 plt.ylabel('Variance')
8
9 plt.subplot(1,2,2)
10 plt.plot(np.arange(1,11), pca.explained_variance_ratio_, 'o-')
11 plt.xlabel('PC number')
12 _ = plt.ylabel('Variance explained')
```



In [17]:

```
1 # Let's look at some of the transformed data, directly.
2
3 plt.figure(figsize=(15, 5))
4
5 # Plot the top two (largest variance) PCs against each other
6 plt.subplot(1,2,1)
7 plt.scatter(X_pca[:,0], X_pca[:,1], c=t)
8 plt.xlabel('PC 1: var=%0.2f' % (np.var(X_pca[:,0])), fontsize=16)
9 plt.ylabel('PC 2: var=%0.2f' % (np.var(X_pca[:,1])), fontsize=16)
10
11 # Plot the bottom two (smallest variance) PCs against each other
12 plt.subplot(1,2,2)
13 plt.scatter(X_pca[:,8], X_pca[:,9], c=t)
14 plt.xlabel('PC 9: var=%0.2f' % (np.var(X_pca[:,8])), fontsize=16)
15 plt.ylabel('PC 10: var=%0.2f' % (np.var(X_pca[:,9])), fontsize=16)
16
17 plt.subplots_adjust(wspace=0.3)
18 plt.show()
```



Ok. There's not a lot to see there. A couple observations:

- The variance of PC1 is notably higher compared to PC2, but variance appears to diminish gradually after that.
- This joint distribution does not look jointly Gaussian (a multivariate Gaussian distribution), at least for the top two PCs.

Let's look at another preprocessing technique...

An alternative (or additional) preprocessing step to scaling is normalization.

Scaling: For all the samples values of a given feature, compute scaling parameters and apply them to all those sample values.

Normalization: For an individual sample, scale all the feature values. Different samples are scaled by different amount. E.g., scale such that the sample has unit norm (lies on the unit circle / sphere / hypersphere).

```
In [18]: 1  # Let's use unit normalization as a preprocessing step.
          2
          3  from sklearn.preprocessing import Normalizer
          4
          5  # Calculate the normalization parameters
          6  normalized = Normalizer()
          7  normalized.fit(X)
          8
          9  # Apply the normalization to our original data
         10  X_norm = normalized.transform(X)
         11
         12  # print(X_norm)
         13  X_norm.shape
```

Out[18]: (30000, 91)



```
In [19]: 1  # Now we'll try PCA again
          2
          3  # Use PCA to compute the transformation parameters for the normalized data
          4  pca = PCA(n_components=10)
          5  pca.fit(X_norm)
          6
          7  # Apply the transformation to our normalized data
          8  X_pca = pca.transform(X_norm)
```

```
In [20]: 1  X.shape, X_pca.shape
```

Out[20]: ((30000, 91), (30000, 10))



```
In [21]: 1 # Let's look at the variance (and variance explained) for each component
2
3 plt.figure(figsize=(15, 5))
4 plt.subplot(1,2,1)
5 plt.plot(np.arange(1,11), pca.explained_variance_, 'o-')
6 plt.xlabel('PC number')
7 plt.ylabel('Variance')
8
9 plt.subplot(1,2,2)
10 plt.plot(np.arange(1,11), pca.explained_variance_ratio_, 'o-')
11 plt.xlabel('PC number')
12 _ = plt.ylabel('Variance explained')
13
14 print(pca.explained_variance_)
15 print(pca.explained_variance_ratio_)
```

```
[0.18087185 0.01926118 0.00697569 0.00531691 0.00441784 0.00388976
 0.00349163 0.00284964 0.0026358  0.00161054]
[0.7717392  0.08218308 0.02976367 0.02268605 0.01884993 0.01659673
 0.01489799 0.01215878 0.01124636 0.00687181]
```



In [22]:

```
1 # Let's look at some of the transformed data, directly.
2
3 plt.figure(figsize=(15, 5))
4
5 # Plot the top two (largest variance) PCs against each other
6 plt.subplot(1,2,1)
7 plt.scatter(X_pca[:,0], X_pca[:,1], c=t)
8 plt.xlabel('PC 1: var=%0.2f' % (np.var(X_pca[:,0])), fontsize=16)
9 plt.ylabel('PC 2: var=%0.2f' % (np.var(X_pca[:,1])), fontsize=16)
10
11 # Plot the bottom two (smallest variance) PCs against each other
12 plt.subplot(1,2,2)
13 plt.scatter(X_pca[:,8], X_pca[:,9], c=t)
14 plt.xlabel('PC 9: var=%0.2f' % (np.var(X_pca[:,8])), fontsize=16)
15 plt.ylabel('PC 10: var=%0.2f' % (np.var(X_pca[:,9])), fontsize=16)
16
17 plt.subplots_adjust(wspace=0.3)
18 plt.show()
```



With normalization, the data variance of the data is well explained by just two components. A model built with only those two components may performance just as well (or better!) than one built with all of the original 91 features.

Take home points

- PCA can (sometimes):
 - Allow us to reduce the number of feature dimensions in our model, reducing the risk of overfitting
 - Allow us to make meaningful data visualization from the top 1 to 3 components
- Scaling and normalization can (sometimes):
 - Allow for a smaller number of PCA dimensions to explain a larger amount of the overall data variation (and thus allow us to use even fewer dimensions in our model.
 - E.g., with normalization, most of the Taiwan data is explained by just the top 2 PCs. Without normalization we may have needed upwards of 8 PCs.

Next Time

- Clustering

