


Introduction to Data Science

19-Clustering



This Lecture

- Clustering

Setup


The obligatory setup code.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import sklearn as sk
        4 import matplotlib.pyplot as plt
        5
        6 from pandas import DataFrame
        7
        8 plt.style.use("ggplot")
        9
       10 %matplotlib inline
       11 %config InlineBackend.figure_format = 'retina'
```

```
In [2]: 1 # function for generating normally distributed data/
        2 def sample_cluster(n, x, y, sigma):
        3     x = np.random.randn(n) * sigma + x;
        4     y = np.random.randn(n) * sigma + y;
        5     return np.array([x, y]).T
        6
```

Synthetic Clustering Example

```
In [3]: 1 np.random.seed(1234)
2 n = 150
3
4 # create three 'sets' of 150 random data
5 c1 = sample_cluster(n, 0, 0.5, 0.15)
6 c2 = sample_cluster(n, 0, 1.2, 0.1)
7 c3 = sample_cluster(n, 0.5, 1, 0.2)
8
9 # put all of that data into one and plot it
10 data = np.concatenate((c1, c2, c3))
11 plt.plot(data[:,0], data[:,1], 'k.')
12 plt.show()
```



Clustering

Simple idea:

- Choose *[Math Processing Error]* (decide how many clusters you *think* there ought to be)
- Partition the data into *[Math Processing Error]* disjoint clusters such that *[Math Processing Error]* is minimized.
 - $W(C)$ is some measure of within-cluster variation
 - Common choice for W is the sum of Euclidean distances between pairs of points:

[Math Processing Error]

Solution Approaches

- Brute force: try every possible partitions of the data
 - Advantage: globally optimal
 - Disadvantage: *[Math Processing Error]* different partitions (given *[Math Processing Error]* data points)
- K-Means algorithm: iterative improvement
 - Advantage: efficient
 - Disadvantage: locally optimal, result depends on initialization

K-Means Algorithm

- Various initialization schemes

- One way is to assign each point a cluster identity at random
- Another is to choose K points to serve as initial cluster centers at random
- Iterate until cluster assignments stop changing:
 - Compute the cluster *centroids*. The centroid is the mean point (vector of means of all features) for all points in the cluster
 - Re-assign each data point to the cluster associated with the nearest centroid

This algorithm is ridiculously easy to code up.

```
In [5]: 1 # a simple k-means implementation
2 # input: integer K, ndarray data
3 def KMeans(K, data):
4     n = data.shape[0]
5     # initialize data into random clusters
6     c = np.random.randint(low=0, high=K, size=n)
7     C = [c]
8     changed = True
9     centroid = lambda k: np.mean(np.array([d for c,d in zip(c,data) if c == k]), axis=0)
10    while changed:
11        centroids = np.array([centroid(k) for k in range(K)])
12        dists = [[np.linalg.norm(d - centroids[k]) for k in range(K)] for d in data]
13        newc = np.argmin(np.array(dists), axis=1)
14        changed = np.any(c != newc)
15        C.append(newc)
16        c = newc
17    return C
```


Let's give it a spin:

```
In [6]: 1 C = KMeans(3, data)
2 print("Iterations:", len(C))
```

Iterations: 7


Let's look at the first few iterations:

```
In [7]: 1 from matplotlib.colors import ListedColormap
2 cmap = ListedColormap(['r', 'y', 'b'])
3 plt.figure(figsize=[20,5])
4 for i in range(3):
5     plt.subplot(1,3,i+1)
6     plt.scatter(data[:,0], data[:,1], c=C[i], cmap=cmap)
```



Look at the final four:

```
In [8]: 1 plt.figure(figsize=[20,5])
2 for i in range(3):
3     plt.subplot(1,3,i+1)
4     plt.scatter(data[:,0], data[:,1], c=C[-(3-i)], cmap=cmap)
```



Here's an implementation that picks K points at random and clusters on those points initially:


```
In [4]: 1 # a simple k-means implementation
2 # input: integer K, ndarray data
3 def KMeans2(K, data):
4     n = data.shape[0]
5     # initialize data using normally distributed random centers
6     centroids = data[np.random.randint(low=0, high=n, size=3)]
7     dists = [[np.linalg.norm(d - centroids[k]) for k in range(K)] for d in data]
8     c = np.argmin(np.array(dists), axis=1)
9     C = [c]
10    changed = True
11    centroid = lambda k: np.mean(np.array([d for c,d in zip(c,data) if c == k]), axis=0)
12    while changed:
13        centroids = np.array([centroid(k) for k in range(K)])
14        dists = [[np.linalg.norm(d - centroids[k]) for k in range(K)] for d in data]
15        newc = np.argmin(np.array(dists), axis=1)
16        changed = np.any(c != newc)
17        C.append(newc)
18        c = newc
19    return C
```

```
In [5]: 1 C = KMeans2(3, data)
2 print("Iterations:", len(C))
```

Iterations: 5


First few iterations:

```
In [11]: 1 plt.figure(figsize=[20,5])
2         for i in range(3):
3             plt.subplot(1,3,i+1)
4             plt.scatter(data[:,0], data[:,1], c=C[i], cmap=cmap)
```



Final four:

```
In [12]: 1 plt.figure(figsize=[20,5])
2         for i in range(3):
3             plt.subplot(1,3,i+1)
4             plt.scatter(data[:,0], data[:,1], c=C[-(3-i)], cmap=cmap)
```



Of course, scikit-learn gives us a robust KMeans object.

It behaves just like a classifier, except that you only give it inputs (no class labels).


The scikit-learn KMeans clusterer will actually try several random starts, and pick the "best" result.

Unfortunately, it doesn't give us a way to visualize the steps it went through.

```
In [13]: 1 from sklearn.cluster import KMeans
2         km = KMeans(3)
3         km.fit(data)
```

```
Out[13]: KMeans(n_clusters=3)
```

```
In [13]: 1 plt.scatter(data[:,0], data[:,1], c=km.predict(data), cmap=cmap)
          2 plt.show()
```



Guessing K


I know the data used above was generated from 3 Gaussian blobs, so we've been using K.

In general, though, how many clusters should there be?

How can we guess?


Let's try K=2:

```
In [14]: 1 km2 = KMeans(2) # uses two clusters
2 km2.fit(data)
3 plt.scatter(data[:,0], data[:,1], c=km2.predict(data), cmap=cmap)
4 plt.show()
```




Or K=4:

```
In [20]: 1 cmap = ListedColormap(['r', 'y', 'b', 'm'])
2 km4 = KMeans(4) # uses two clusters
3 km4.fit(data)
4 plt.scatter(data[:,0], data[:,1], c=km4.predict(data), cmap=cmap)
5 plt.show()
```



Or K=10:

```
In [21]: 1 km10 = KMeans(10) # uses ten clusters
2 km10.fit(data)
3 plt.scatter(data[:,0], data[:,1], c=km10.predict(data), cmap='nipy_spectral')
4 plt.show()
```



Hierarchical Clustering

One answer to the problem of "how many clusters":

Try all of them!


Hierarchical algorithms are either top down or bottom up.

The bottom up algorithms (agglomerative) are the most popular.

The result of hierarchical clustering is a *dendrogram* plot, showing how the clusters break down (or build up, if you prefer).

Let's first get a smaller dataset to illustrate.


```
In [22]: 1 c1 = sample_cluster(10, 0, 0.5, 0.05)
2 c2 = sample_cluster(10, 0, 1.2, 0.1)
3 c3 = sample_cluster(10, 0.5, 1, 0.05)
4 data = np.concatenate((c1, c2, c3))
5 plt.scatter(data[:,0], data[:,1])
6 plt.show()
```



While scikit-learn has agglomerative clustering, it isn't capable of showing us a dendrogram.

We'll use SciPy instead for this example.

```
In [18]: 1 from scipy.cluster.hierarchy import dendrogram, ward
2 linkage_array = ward(data)
3 plt.figure(figsize=[20,5])
4 dendrogram(linkage_array)
5 plt.show()
```

